

Objective of this exercise is to access Hadoop hdfs using JAVA API's from our local environment.
Certain changes to be done before accessing HDFS

Changes in Vm-ware

- 1) Change the fs.default.name to hdfs://<IP OF VM>:8020 in /etc/hadoop-0.20/conf/core-site.xml
- 2) In mapred-site.xml mapred.job.tracker to <IP OF VM>:8021
- 3) In /etc/hosts make sure that there is no entry of 127.0.0.1 and there is a hostname for your VM IP
IP of your VM can be found by typing "ifconfig" in the terminal of your VM ware.

Changes in Local Host Environment

In running from windows C:\Windows\System32\drivers\etc\hosts, add the mapping
192.168.91.128 cloudera-vm
(ip of your vm)

To make your changes effective in hadoop you will have to restart namenode, data node and jobtracker, use command

```
sudo service hadoop-hdfs-namenode restart
sudo service hadoop-hdfs-datanode restart
sudo service hadoop-0.20-mapreduce-jobtracker restart
sudo service hadoop-0.20-mapreduce-tasktracker restart
```

After doing this do health check for your hdfs system. Hadoop fsck / -files -blocks

If the Status is / healthy, changes have been applied.

```
import java.io.BufferedInputStream;
import java.io.BufferedReader;
import java.io.File;
import java.io.FileInputStream;
import java.io.IOException;
import java.io.InputStream;
import java.io.InputStreamReader;
import java.util.ArrayList;
```

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FSDataOutputStream;
import org.apache.hadoop.fs.FileStatus;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.FileUtil;
import org.apache.hadoop.fs.Path;
```

```
public class HDFSClient {
```

```
    private static String HDFS_Host_name="hdfs://192.168.91.128/";
```

```
    public void printUsage(){
        System.out.println("Usage: hdfsclient read <hdfs_path>");
        System.out.println("Usage: hdfsclient delete <hdfs_path>");
        System.out.println("Usage: hdfsclient mkdir <hdfs_path>");
        System.out.println("Usage: hdfsclient rename_file <hdfs_path>");
        System.out.println("Usage: hdfsclient add_file <local_path> <hdfs_path>");
        System.out.println("Usage: hdfsclient CopyToLocal <hdfs_path> <local_path>");
        System.out.println("Usage: hdfsclient CopyFromLocal <local_path> <hdfs_path>");
    }
    public void read(String file) throws IOException{
        Path path = new Path(HDFS_Host_name+file);
        Configuration conf = new Configuration();
```

Files To be imported

HDFS path name and
executable methods displayed
in printUsage()

```

conf.addResource(new Path("/usr/lib/hadoop-0.20-mapreduce/conf/core-site.xml"));
FileSystem fs = FileSystem.get(path.toUri(),conf);
if(!fs.exists(path)){
    System.out.println("File "+ file +" does not exists!");
    return;
}
BufferedReader br = new BufferedReader(new InputStreamReader(fs.open(path)));
String s = br.readLine();
while(s!=null){
    System.out.println(s);
    s = br.readLine();
}
br.close();
fs.close();
}

public void delete(String file) throws IOException{
    Path path = new Path(HDFS_Host_name+file);
    Configuration conf = new Configuration();
    conf.addResource(new Path("/usr/lib/hadoop-0.20-mapreduce/conf/core-site.xml"));
    FileSystem fs = FileSystem.get(path.toUri(),conf);
    if(!fs.exists(path)){
        System.out.println("File "+ file +" does not exists!");
        return;
    }
    fs.delete(path, true);
    fs.close();
}

public void mkdir(String file) throws IOException{
    Path path = new Path(HDFS_Host_name+file);
    Configuration conf = new Configuration();
    FileSystem fs = FileSystem.get(path.toUri(),conf);
    conf.addResource(new Path("/usr/lib/hadoop-0.20-mapreduce/conf/core-site.xml"));
    if(fs.exists(path)){
        System.out.println("Dir. "+ file +" already exists");
        return;
    }
    fs.mkdirs(path);
    fs.close();
}

public void rename_file(String from_file,String to_file) throws IOException{
    Path from_path = new Path(HDFS_Host_name+from_file);
    Path to_path = new Path(HDFS_Host_name+to_file);
    Configuration conf = new Configuration();
    conf.addResource(new Path("/usr/lib/hadoop-0.20-mapreduce/conf/core-site.xml"));
    FileSystem fs = FileSystem.get(from_path.toUri(),conf);
    if(!fs.exists(from_path)){
        System.out.println("File "+ from_file +" does not exists");
        return;
    }
    if(fs.exists(to_path)){
        System.out.println("File "+ to_file +" already exists");
        return;
    }
    fs.rename(from_path, to_path);
    fs.close();
}

public void add_file(String source,String dest) throws IOException{
    Path path = new Path(HDFS_Host_name+dest);
    Configuration conf = new Configuration();
    conf.addResource(new Path("/usr/lib/hadoop-0.20-mapreduce/conf/core-site.xml"));
    conf.addResource(new Path("/usr/lib/hadoop-0.20-mapreduce/conf/hdfs-site.xml"));
    FileSystem fs = FileSystem.get(path.toUri(),conf);
    if(fs.exists(path)){
        System.out.println("File "+ dest +" already exists");
        return;
    }
}

```

File System provides various methods to interact with HDFS

```

FSDDataOutputStream out = fs.create(path);
InputStream in= new BufferedInputStream(new FileInputStream(new File(source)));
byte[] b =new byte[1024];
int numbytes=0;
while((numbytes=in.read(b))>0){
    out.write(b, 0, numbytes);
}
in.close();
out.close();
fs.close();
}

```

InputStream reads the file as a stream of bytes. These bytes stored in array temporarily before being written into FSDDataOutputStream which writes these stream of bytes in HDFS specified by path.

```

}
public void CopyToLocal(String source,String dest) throws IOException{
    Path path1 = new Path(HDFS_Host_name+source);
    Configuration conf = new Configuration();
    conf.addResource(new Path("/usr/lib/hadoop-0.20-mapreduce/conf/core-site.xml"));
    FileSystem fs1 = FileSystem.get(path1.toUri(),conf);
    if(!fs1.exists(path1)){
        System.out.println("File "+ source +" does not exists");
        return;
    }
    Path path2 = new Path(dest);
    FileSystem fs2 = FileSystem.get(path2.toUri(),conf);
    if(fs2.exists(path2)){
        System.out.println("File "+ dest +" already exists");
        return;
    }
    FileUtil.copy(fs1, path1, new File(dest), false, conf);
}
}
public void CopyFromLocal(String source, String dest) throws IOException {

```

In this method , copy action takes place using predefined classes and methods of FileUtil and FileSystem.

```

    Path destPath = new Path(HDFS_Host_name + dest);
    Configuration conf = new Configuration();
    conf.addResource(new Path("/usr/lib/hadoop-0.20-mapreduce/conf/core-site.xml"));
    FileSystem hdfsfileSystem = FileSystem.get(destPath.toUri(), conf);

    Path srcPath = new Path(source);
    FileSystem localfileSystem = FileSystem.get(srcPath.toUri(),conf);
    try {

        ArrayList<Path> srcPaths = getSourcePaths(localfileSystem, srcPath);
        Path[] sourcePaths = srcPaths.toArray(new Path[srcPaths.size()]);
        FileUtil.copy(localfileSystem, sourcePaths, hdfsfileSystem, destPath, false, false, conf);
        System.out.println("File " + srcPaths + "copied to " + dest);
    } catch (Exception e) {
        System.err.println("Exception caught! ." + e);
        e.printStackTrace();
        System.exit(1);
    } finally {
        localfileSystem.close();
        hdfsfileSystem.close();
    }
}

```

```

private ArrayList<Path> getSourcePaths(FileSystem localfileSystem, Path srcPath) throws Exception{
    ArrayList<Path> sourcePaths = new ArrayList<Path>();
    for (FileStatus file : localfileSystem.listStatus(srcPath)){
        if(file.isDir()){
            sourcePaths.addAll(getSourcePaths(localfileSystem,file.getPath()));
        }
        else
            sourcePaths.add(file.getPath());
    }
    return sourcePaths;
}
}

```

FileStatus lists all the paths of files and directories if specified FileSystem is directory. A loop is run to make sure all the paths of files in this filesystem are written in ArrayList<Path> source paths

Main Class

```
import java.io.IOException;
```

```
public class Main {
    public static void main(String[] args) throws IOException {
        HDFSClient client = new HDFSClient();

        if (args.length < 1) {
            client.printUsage();
            System.exit(1);
        }

        if (args[0].equals("add")) {
            if (args.length < 3) {
                System.out.println("Usage: hdfsclient add_file <local_path> <hdfs_path>");
                System.exit(1);
            }
            client.add_file(args[1], args[2]);
        } else if (args[0].equals("read")) {
            if (args.length < 2) {
                System.out.println("Usage: hdfsclient read <hdfs_path>");
                System.exit(1);
            }
            client.read(args[1]);
        } else if (args[0].equals("delete")) {
            if (args.length < 2) {
                System.out.println("Usage: hdfsclient delete <hdfs_path>");
                System.exit(1);
            }
            client.delete(args[1]);
        } else if (args[0].equals("mkdir")) {
            if (args.length < 2) {
                System.out.println("Usage: hdfsclient mkdir <hdfs_path>");
                System.exit(1);
            }
            client.mkdir(args[1]);
        } else if (args[0].equals("CopyFromLocal")) {
            if (args.length < 3) {
                System.out
                    .println("Usage: hdfsclient copyfromlocal <from_local_path>
<to_hdfs_path>");
                System.exit(1);
            }
            client.CopyFromLocal(args[1], args[2]);
        } else if (args[0].equals("rename")) {
            if (args.length < 3) {
                System.out
```

```
                .println("Usage: hdfsclient rename <old_hdfs_path> <new_hdfs_path>");
                System.exit(1);
            }

            client.rename_file(args[1], args[2]);
        } else if (args[0].equals("CopyToLocal")) {
            if (args.length < 3) {
                System.out
                    .println("Usage: hdfsclient copytolocal <from_hdfs_path>
<to_local_path>");
                System.exit(1);
            }

            client.CopyToLocal(args[1], args[2]);
        } else {
            client.printUsage();
            System.exit(1);
        }
        System.out.println("Done!");
    }
}
```