

Assignment 4

In this Exercise, we have included DistributedCache class to import a file containing words to be skipped while calculating their word count. Also counters are used to report corresponding type of words through the use of ENUM.

EnhancedWordCount Class

```
import java.io.*;
import java.util.*;
```

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.filecache.DistributedCache;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
```

```
enum NatureofWords { STARTS_WITH_DIGIT, STARTS_WITH_LETTER, ALL }
```

Enum to keep track counter of words starting with a digit/letter and all words counted

```
public class EnhancedWordCount extends Configured {
    public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
```

```
        private boolean caseSensitive = true;
        private Set<String> patternsToSkip = new HashSet<String>();
```

The Default value of casesensitive Boolean. If not specifically specified through command line, this program differentiates between upper case and lower case letters .

```
        private long numRecords = 0;
        private String inputFile;
        private BufferedReader fis;
```

```
        public void configure(JobConf job) {
            caseSensitive = job.getBoolean("wordcount.case.sensitive", true);
            inputFile = job.get("map.input.file");
```

```
            if (job.getBoolean("wordcount.skip.patterns", false)) {
                Path[] patternsFiles = new Path[0];
                try {
                    patternsFiles = DistributedCache.getLocalCacheFiles(job);
                } catch (IOException ioe) {
                    System.err.println("Caught exception getting cached files: " + ioe.toString());
                }
                for (Path patternsFile : patternsFiles) {
                    parseSkipFile(patternsFile);
                }
            }
        }
```

We will be mentioning the Boolean value of wordcount.skip.patterns in command line through string "-skip".

```
        private void parseSkipFile(Path patternsFile) {
            try {
                fis = new BufferedReader(new FileReader(patternsFile.toString()));
```

Paths entered in DistributedCache are taken out. These paths are fed in parseSkipFile(). This function reads individual words from the files referred by extracted paths and added to patternsToSkip array.

```
String pattern = null;
while ((pattern = fis.readLine()) != null) {
    patternsToSkip.add(pattern);}
} catch (IOException ioe) {
    System.err.println("Caught exception parsing the cached file '"+patternsFile+"' "+ ioe.toString());
}
}}
```

```
public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter)
throws IOException {
    String line = (caseSensitive) ? value.toString() : value.toString().toLowerCase();
    StringTokenizer tokenizer = new StringTokenizer(line);
    while (tokenizer.hasMoreTokens()) {
        String token = tokenizer.nextToken();
        if (patternsToSkip.contains(token))
            System.out.println("Word Skipped");
        else{
            word.set(token);
            output.collect(word, one);
        }
    }
    //reporter.incrCounter(Counters.INPUT_WORDS, 1);
}}
```

- 1) If Boolean "caseSensitive" is true, string value is taken as it is, if false, string value is taken as lower case.
- 2) If words extracted are present in patternsToSkip array, they are not sent to output.collect

```
public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output,
Reporter reporter) throws IOException {
        int sum = 0;

        String token = key.toString();
        if (StringUtils.startsWithDigit(token)){
            reporter.incrCounter(NatureofWords.STARTS_WITH_DIGIT, 1);
        }

        else if (StringUtils.startsWithLetter(token)){
            reporter.incrCounter(NatureofWords.STARTS_WITH_LETTER, 1);
        }

        reporter.incrCounter(NatureofWords.ALL, 1);

        while (values.hasNext()) {
            sum += values.next().get();
        }

        output.collect(key, new IntWritable(sum));
    }
}
```

Counters for defined Enum are increased from output of StringUtils class defined.

```
public static void main (String[] args) throws Exception {
    JobConf conf = new JobConf(EnhancedWordCount.class);
    conf.setJobName("enhancedwordcount");
}
```

```

conf.setOutputKeyClass(Text.class);
conf.setOutputValueClass(IntWritable.class);
conf.setMapperClass(Map.class);
conf.setCombinerClass(Reduce.class);
conf.setReducerClass(Reduce.class);

conf.setInputFormat(TextInputFormat.class);
conf.setOutputFormat(TextOutputFormat.class);

List<String> other_args = new ArrayList<String>();
for (int i=0; i < args.length; ++i) {
    if ("-skip".equals(args[i])) {
        DistributedCache.addCacheFile(new Path(args[++i]).toUri(), conf);
        conf.setBoolean("wordcount.skip.patterns", true);
    } else if ("-case".equals(args[i])) {
        conf.setBoolean("wordcount.case.sensitive", false );
    }
    other_args.add(args[i]);
}

FileInputFormat.setInputPaths(conf, new Path(other_args.get(0)));
FileOutputFormat.setOutputPath(conf, new Path(other_args.get(1)));

JobClient.runJob(conf);
}

```

The above for loop takes inputs as command line arguments

- 1) If string “-skip” is encountered, next string of the command line is taken as address of the file which contains words to be skipped. The file associated with this string is added to DistributedCache class.
- 2) If string “-case” is encountered, the Boolean of wordcount.case.sensitive is set as false and program doesnot differentiates between upper and lower case letters.

StringUtil Class

```

public class StringUtils {
    public static boolean startsWithDigit(String s){
        if( s == null || s.length() == 0 )
            return false;

        return Character.isDigit(s.charAt(0));
    }

    public static boolean startsWithLetter(String s){
        if( s == null || s.length() == 0 )

```

This Class defines two functions startsWithDigit() and StartsWithLetter() which return Boolean.

```
return false;
```

```
return Character.isLetter(s.charAt(0));
```

```
}}
```

Running the Program:

```
File Edit View Search Terminal Help
cloudera@cloudera-vm:~/Desktop$
cloudera@cloudera-vm:~/Desktop$
cloudera@cloudera-vm:~/Desktop$
cloudera@cloudera-vm:~/Desktop$ hadoop jar ewc.jar EnhancedWordCount /user/cloudera/UN.txt /user/cloudera/skip.txt -case -skip
File with words to skip string for not checking case sensitivity. String for skipping words

Applications Places System
cloudera@cloudera-vm: ~/Desktop
File Edit View Search Terminal Help
l for the same.
14/01/28 02:07:23 INFO mapred.FileInputFormat: Total input paths to process : 1
14/01/28 02:07:24 INFO mapred.JobClient: Running job: job_201401272358_0010
14/01/28 02:07:25 INFO mapred.JobClient: map 0% reduce 0%
14/01/28 02:08:47 INFO mapred.JobClient: map 100% reduce 0%
14/01/28 02:08:57 INFO mapred.JobClient: map 100% reduce 100%
14/01/28 02:08:58 INFO mapred.JobClient: Job complete: job_201401272358_0010
14/01/28 02:08:58 INFO mapred.JobClient: Counters: 26
14/01/28 02:08:58 INFO mapred.JobClient: Job Counters
14/01/28 02:08:58 INFO mapred.JobClient: Launched reduce tasks=1
14/01/28 02:08:58 INFO mapred.JobClient: SLOTS_MILLIS_MAPS=149974
14/01/28 02:08:58 INFO mapred.JobClient: Total time spent by all reduces waiting after reserving slots (ms)=0
14/01/28 02:08:58 INFO mapred.JobClient: Total time spent by all maps waiting after reserving slots (ms)=0
14/01/28 02:08:58 INFO mapred.JobClient: Launched map tasks=2
14/01/28 02:08:58 INFO mapred.JobClient: Data-local map tasks=2
14/01/28 02:08:58 INFO mapred.JobClient: SLOTS_MILLIS_REDUCE=10204
14/01/28 02:08:58 INFO mapred.JobClient: FilesystemCounters
14/01/28 02:08:58 INFO mapred.JobClient: FILE_BYTES_READ=3971
14/01/28 02:08:58 INFO mapred.JobClient: HDFS_BYTES_READ=4765
14/01/28 02:08:58 INFO mapred.JobClient: FILE_BYTES_WRITTEN=171846
14/01/28 02:08:58 INFO mapred.JobClient: HDFS_BYTES_WRITTEN=2603
14/01/28 02:08:58 INFO mapred.JobClient: NatureOfWords
14/01/28 02:08:58 INFO mapred.JobClient: STARTS_WITH_DIGIT=12
14/01/28 02:08:58 INFO mapred.JobClient: STARTS_WITH_LETTER=243
14/01/28 02:08:58 INFO mapred.JobClient: ALL=260
14/01/28 02:08:58 INFO mapred.JobClient: Map-Reduce Framework
14/01/28 02:08:58 INFO mapred.JobClient: Reduce input groups=260
14/01/28 02:08:58 INFO mapred.JobClient: Combine output records=286
14/01/28 02:08:58 INFO mapred.JobClient: Map input records=4
14/01/28 02:08:58 INFO mapred.JobClient: Reduce shuffle bytes=3977
14/01/28 02:08:58 INFO mapred.JobClient: Reduce output records=260
14/01/28 02:08:58 INFO mapred.JobClient: Spilled Records=572
14/01/28 02:08:58 INFO mapred.JobClient: Map output bytes=4368
14/01/28 02:08:58 INFO mapred.JobClient: Map input bytes=3053
14/01/28 02:08:58 INFO mapred.JobClient: Combine input records=390
14/01/28 02:08:58 INFO mapred.JobClient: Map output records=390
14/01/28 02:08:58 INFO mapred.JobClient: SPLIT_RAW_BYTES=184
14/01/28 02:08:58 INFO mapred.JobClient: Reduce input records=286
cloudera@cloudera-vm:~/Desktop$
```

If you do want upper and lowercase letters to be counted separately, do not type “-case” in command line. Similarly do not include “-skip” and corresponding file path if you do not want words to be skipped.