

Assignment5

In this assignment of **Relational Join**, we will join two tables. These two tables are present in two files named *id.txt* and *trips.txt*.

id.txt

This file has id wise data for employees company.

```
cloudera@cloudera-vm:~/Desktop$ hadoop fs -cat /user/cloudera/rel_join/id.txt
101    aaa    executive
102    bbb    manager
104    hhh    manager
106    ccc    trainee
109    hhh    trainee
103    ddd    manager
105    bbb    executive
107    eee    trainee
110    fff    vice-president
108    ggg    president
cloudera@cloudera-vm:~/Desktop$
```

trips.txt

This file has trip details for each employee.

```
cloudera@cloudera-vm:~/Desktop$ hadoop fs -cat /user/cloudera/rel_join/trips.txt
101    pune    1
101    hyd     2
102    pune    2
102    hyd     3
102    bang    4
104    pune    2
104    hyd     4
104    bang    5
106    pune    1
109    pune    1
103    pune    2
103    hyd     3
103    bang    5
105    pune    2
105    hyd     4
107    pune    2
110    pune    2
110    hyd     3
110    bang    5
110    del     2
108    pune    2
108    hyd     3
108    bang    4
108    del     1
108    chen    1
cloudera@cloudera-vm:~/Desktop$
```

Objective: We have to join these two data files based on employee id.

Solution:

In this case, we will be giving input to our map reduce program in the form of two files. Certain new commands will be introduced in the *run.class*. Also we will be writing two map classes one for each file to keep the track of our data.

Mapper class for "id.txt"(it's a tab delimited file)

```
import java.io.IOException;
import java.util.StringTokenizer;
```

```
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
```

```
public class map_id extends MapReduceBase implements
Mapper<LongWritable,Text,DoubleWritable,Text> {
    private final String id_file = "zzzzzzzzzz";
    public void map(LongWritable key, Text value,
OutputCollector<DoubleWritable,Text>output,Reporter reporter) throws IOException {
        String line=value.toString();
        String outline=id_file;
        double id=0;
        StringTokenizer token = new StringTokenizer(line);
        if ( token.hasMoreTokens()){
            id = Double.parseDouble(token.nextToken());
        }
        while (token.hasMoreTokens() ){
            outline = outline + '\t' +token.nextToken();
        }
        output.collect(new DoubleWritable(id),new Text(outline));
    }
}
```

Counter added to identify data from "id.txt"

First token is converted to double as employee id which will serve as a key to be sent to reducer program.

- 1) Remaining tokens are added to token "zzzzzzzzzz" for identification during reduce phase.

Reason for adding our 'identification' token.

Output from our mapper function will be sent to reducer function. So all the values with same key from both files will be sent to reducer function. We have added our token so that we can identify values from "id.txt" so that these values can be added to values from "trips.txt" and hence final joined values can be sent to output collector.

Mapper class for "trips.txt"

```
import java.io.IOException;
import java.util.StringTokenizer;
```

```
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
```

```
public class map_trips extends MapReduceBase implements
Mapper<LongWritable,Text,DoubleWritable,Text> {
```

```
    public void map(LongWritable key, Text value,
OutputCollector<DoubleWritable,Text>output,Reporter reporter) throws IOException {
```

```
        String line=value.toString();
        String outline2=null;
        int counter =0;
        double id=0;
        StringTokenizer token = new StringTokenizer(line);

        if ( token.hasMoreTokens()){
            id = Double.parseDouble(token.nextToken());
```

```
        }
        while (token.hasMoreTokens() ){
            if(counter == 1)
                outline2 = outline2 + ('\t'+token.nextToken());
            else{
                outline2 = token.nextToken();
                counter =1;
            }
        }
```

```
        output.collect(new DoubleWritable(id),new Text(outline2));
```

```
    }
}
```

- 1) First token is converted to double as employee id.
- 2) Remaining tokens are added together to form a string and sent as value to reduce phase.

. 'counter' keeps count if we have taken out second token or not. Once second token is taken out,through counter remaining tokens are added to previous string with tab space between them.

Reducer Class

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.StringTokenizer;

import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class reduce extends MapReduceBase implements Reducer<DoubleWritable,Text,DoubleWritable,Text> {
    public void reduce (DoubleWritable key, Iterator<Text> values, OutputCollector<DoubleWritable,Text>
    output, Reporter reporter) throws IOException {
```

```
        ArrayList<String> arr = new ArrayList<String>();
        arr.clear();
        String line = null;
        String base = null;
        int counter = 0;
        int first = 0;
        int i = 0;

        while(values.hasNext()) {
            line = values.next().toString();
            StringTokenizer token = new StringTokenizer(line);
            if(token.countTokens() == 3){
                token.nextToken();
                while (token.hasMoreTokens()){
                    if(first == 0){
                        base = token.nextToken();
                        first = 1;
                    }
                    else
                        base = base + '\t' + token.nextToken();
                }
            }
            else {
                arr.add(line);
                counter++;
            }
        }
        while(i < counter){
            line = base + " " + '\t' + arr.get(i);
            output.collect(key, new Text(line));
            i++;
        }
    }
}
```

First token from "id.txt" is not included because it's our added token for identification. 'first' keeps count if we have taken out second token or not. Once second token is taken out, through first remaining tokens are added to previous string with tab space between them.

- 1) If first value after tokenizing, has 3 tokens (due to our added token "zzzzzz", it is identified as value from id.txt.
- 2) Then first added token is taken out and rest are joined to form base string.
- 3) If number of tokens are not 3, (value from trips.txt) the string is sent into an array list for future reference.

Finally all the strings from array list are first added to base string to be sent as value to output collector with key as employee id giving the final required table.

Runner Class

```
import java.io.IOException;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.TextOutputFormat;
import org.apache.hadoop.mapred.lib.MultipleInputs;
import org.apache.hadoop.mapred.FileInputFormat;

public class run {
    public static void main(String[] args) throws IOException{
        JobConf conf = new JobConf(run.class);
        conf.setJobName("Relational Join");

        conf.setOutputKeyClass(DoubleWritable.class);
        conf.setOutputValueClass(Text.class);

        conf.setReducerClass(reduce.class);

        conf.setOutputFormat(TextOutputFormat.class);

        MultipleInputs.addInputPath(conf, new Path(args[0]), TextInputFormat.class, map_id.class);
        MultipleInputs.addInputPath(conf, new Path(args[1]), TextInputFormat.class, map_trips.class);
        FileOutputFormat.setOutputPath(conf, new Path(args[2]));

        JobClient.runJob(conf);
    }
}
```

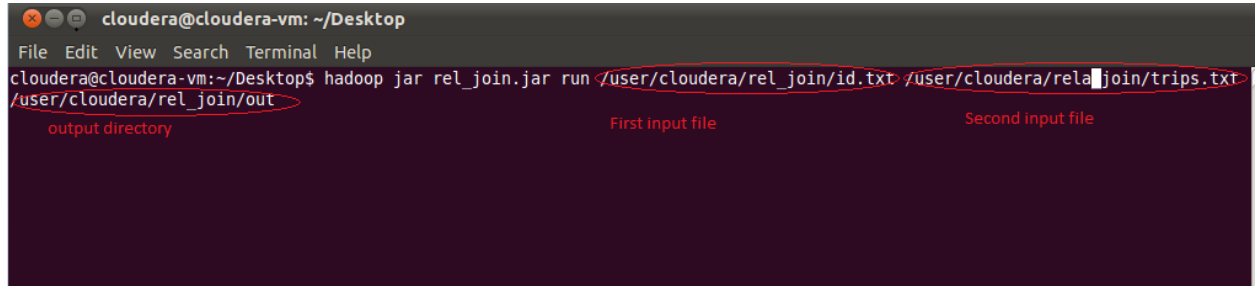
Runner Class

- 1) Initialise a JobConf class mentioning runner class in the constructor.
- 2) Set Output Key and Output Value Class. In our case we have output key as double(employee id) and output value as Text(employee's name, designation, place and roundtrips).
- 3) Reducer class is set. Mapper class is set individually for both inputs.
- 4) Output Format are set. In our case our input data and output result both are in Text Format. MultipleInputs class adds multiple inputs with different mapper classes.
- 5) Paths are mentioned through command line arguments.

MultipleInputs class to add multiple inputs to put map reduce program with arguments as (Jobconf,path,InputFormatclass,mapper class)

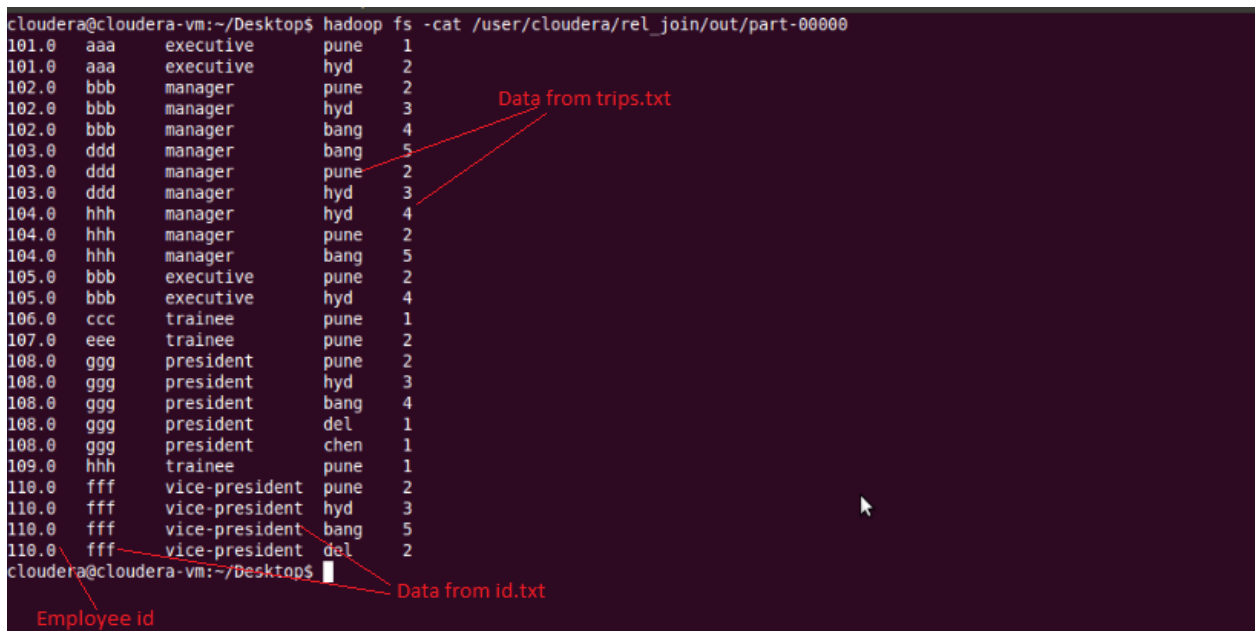
Running the program :

- 1) Make sure both files "id.txt" and "trips.txt" are present in your hdfs system.
- 2) Run the following command .



```
cloudera@cloudera-vm: ~/Desktop
File Edit View Search Terminal Help
cloudera@cloudera-vm:~/Desktop$ hadoop jar rel_join.jar run /user/cloudera/rel_join/id.txt /user/cloudera/rel_join/trips.txt /user/cloudera/rel_join/out
output directory                                First input file                                Second input file
```

- 3) Final output



```
cloudera@cloudera-vm:~/Desktop$ hadoop fs -cat /user/cloudera/rel_join/out/part-00000
101.0 aaa executive pune 1
101.0 aaa executive hyd 2
102.0 bbb manager pune 2
102.0 bbb manager hyd 3
102.0 bbb manager bang 4
103.0 ddd manager bang 5
103.0 ddd manager pune 2
103.0 ddd manager hyd 3
104.0 hhh manager hyd 4
104.0 hhh manager pune 2
104.0 hhh manager bang 5
105.0 bbb executive pune 2
105.0 bbb executive hyd 4
106.0 ccc trainee pune 1
107.0 eee trainee pune 2
108.0 ggg president pune 2
108.0 ggg president hyd 3
108.0 ggg president bang 4
108.0 ggg president del 1
108.0 ggg president chen 1
109.0 hhh trainee pune 1
110.0 fff vice-president pune 2
110.0 fff vice-president hyd 3
110.0 fff vice-president bang 5
110.0 fff vice-president del 2
cloudera@cloudera-vm:~/Desktop$
```

Employee id

Data from id.txt

Data from trips.txt